

# Trustworthy Federated Graph Neural Networks: Unifying Fairness, Privacy, Robustness, and Explainability for High-Stakes Risk Detection

Vinh Dang<sup>1\*</sup>

<sup>1\*</sup>Department of Computer Science, Independent Researcher, Vietnam.

Corresponding author(s). E-mail(s): [dqvinh87@gmail.com](mailto:dqvinh87@gmail.com);

## Abstract

Federated Learning (FL) enables institutions to train shared models over sensitive, decentralised graph data—transactions, credit records, social ties—without exchanging raw data. Yet accuracy alone does not make such systems deployable: recent regulatory and scientific frameworks (the EU AI Act, the NIST AI Risk Management Framework) demand that automated decision systems be simultaneously *fair*, *private*, *robust*, and *transparent*. Existing work addresses these properties in isolation, and no prior method combines fairness, differential privacy, and Byzantine robustness for federated graph neural networks (GNNs). We present **FedFairGNN**, a unified trustworthy federated graph-learning framework built from three components: (i) *Fairness-Sensitive Edge Reweighting* (FSER), which suppresses the biased cross-group message passing responsible for structural discrimination; (ii) *Fairness-Targeted Gradient Decomposition* (FTGD), which achieves differential privacy for the sensitive attribute by privatising only the low-dimensional demographic-parity statistics—making privacy nearly free where standard full-gradient DP-SGD destroys utility; and (iii) *Bi-objective Frank–Wolfe Aggregation* (BFWA) and its Byzantine-robust variant, which enforce a hard fairness budget during aggregation while screening malicious clients, including a new fairness-poisoning threat. We further instrument the framework with predictive uncertainty, GNN attention explanations, a composite trust score, sustainability accounting, and an EU AI Act / NIST RMF compliance mapping. Across five real datasets spanning credit, recidivism, social, and—uniquely—large-scale crypto (Elliptic Bitcoin, 203,769 nodes) domains, FedFairGNN attains the best fairness–utility–privacy–robustness trade-off among ten baselines, including recent fairness-aware federated and graph methods. All tables and figures are regenerated from logged runs; code and experiments are released for full reproducibility. This work supports SDG 9 by advancing trustworthy, infrastructure-grade collaborative AI.

**Keywords:** Federated learning, Graph neural networks, Algorithmic fairness, Differential privacy, Byzantine robustness, Explainable AI, Trust metrics, Regulatory compliance

## 1 Introduction

Consider a consortium of banks that together want to catch money laundering. Fraud is fundamentally *relational*: illicit funds move through chains of accounts, shell companies, and intermediaries, and it is precisely the *structure* of these money flows—not any single transaction in isolation—that betrays criminal intent. Graph Neural Networks (GNNs) have therefore become the tool of choice for financial crime detection, credit scoring, and other high-stakes risk assessments, because their message-passing mechanism learns directly from the web of relationships between users, accounts, and transactions [1–3]. Yet the very banks that would benefit most from a shared model are legally forbidden from pooling their customers’ data: privacy regulations such as the GDPR, and plain commercial confidentiality, keep each institution’s transaction graph locked inside its own walls. Federated Learning (FL) offers a way out. By exchanging model updates instead of raw data, a group of institutions can train a single, more powerful detector without any of them surrendering its records [4].

This picture, however, is dangerously incomplete. A model that is merely *accurate* is no longer fit to be deployed in decisions that shape people’s access to credit, liberty, or financial services. Three converging forces—the European Commission’s Ethics Guidelines for Trustworthy AI, the EU AI Act, and the U.S. NIST AI Risk Management Framework—now demand that automated decision systems be simultaneously *fair*, *privacy-preserving*, *robust*, *transparent*, and *accountable*. These are not optional extras: for a high-risk system under the EU AI Act they are legal prerequisites for going to market. Meeting them is hard even for a centralised model. In the federated, graph-structured setting they collide, and each one becomes markedly harder than in isolation.

### ***Why the federated graph setting is uniquely hard.***

The difficulty is not that we must bolt four independent features onto a classifier; it is that the natural solution to each requirement actively sabotages the others. *Fairness* is confounded by the graph itself: homophilous message passing means that a member of a demographic group who happens to be connected to risky neighbours will be flagged as risky too, so the model learns to launder a protected attribute through graph structure—and the central server, which never sees any data, cannot even measure this bias, let alone correct it. *Privacy* appears to have a textbook answer, differentially private SGD, but that answer fails badly here: adding calibrated noise to every coordinate of the gradient injects perturbation whose magnitude grows with the square root of the model dimension, and on the compact GNNs used for tabular financial graphs this noise simply drowns the signal, collapsing accuracy to little better than chance (we quantify this in Section 5.5). *Robustness* is threatened not only by the classical malfunctioning or malicious client, but—as we show—by a subtler adversary that the

fairness machinery itself invites: a client can transmit a deliberately biased model while *reporting* exemplary fairness statistics, thereby hijacking any aggregator that rewards fairness. And *transparency* is undermined by the opacity of relational models, exactly when regulation requires that each adverse decision be explainable and auditable.

### *The gap.*

The research literature has, understandably, attacked these problems one at a time. Fair-GNN methods deliver impressive demographic parity but assume all data sits in one place [5–7]. Fair federated learning has matured for tabular and image data but rarely touches graphs [8, 9]. Byzantine-robust aggregation hardens FL against poisoning but is oblivious to fairness [10, 11]. Differentially private FL protects data but, as noted, at a utility cost that is prohibitive for graph models. To the best of our knowledge, *no* existing method delivers fairness, differential privacy, and Byzantine robustness together for federated GNNs, and none has been validated on a large-scale cryptocurrency graph—arguably the most consequential modern fraud domain. Closing this gap is the goal of this paper.

### *Our approach.*

We present **FedFairGNN**, a single framework in which fairness, privacy, robustness, and transparency are designed to reinforce rather than undermine one another. Its architecture (Fig. 1) intervenes at the three natural control points of federated graph training—inside message passing, at the local gradient, and at server aggregation—and wraps the result in an instrumentation layer that measures, explains, and audits trustworthiness. The key technical insight, and the reason privacy becomes almost free, is a change of perspective on *where* the sensitive attribute actually influences learning: it does so only through a two-dimensional summary statistic, so we can protect that statistic directly instead of blindly noising the entire gradient.

Concretely, our contributions are:

1. **Fairness-Sensitive Edge Reweighting (FSER)**, a learnable attention correction that detects and suppresses the biased cross-group edges responsible for structural discrimination, mitigating unfairness at its source inside the GNN rather than patching it after the fact (Section 4.1).
2. **Fairness-Targeted Gradient Decomposition (FTGD)**. We observe that the sensitive attribute enters the model *only* through the demographic-parity statistics, and exploit this to obtain  $(\epsilon, \delta)$ -differential privacy by privatising two bounded-sensitivity scalars instead of the full gradient. This makes strong privacy essentially free—precisely where standard DP-FedAvg collapses utility—and we prove the guarantee with a self-contained Rényi-DP accountant (Section 4.2).
3. **Bi-objective Frank–Wolfe Aggregation (BFWA)** and a **Byzantine-robust variant**. The server solves a constrained optimisation that enforces a hard, operator-chosen fairness budget during aggregation, while distance-based screening neutralises poisoning—including the fairness-poisoning attack of [12], which we reproduce and defend against (Sections 4.3, 5.7).

4. **A trustworthiness instrumentation layer** that turns abstract principles into measured, auditable artifacts: predictive uncertainty and calibration, GNN attention explanations with fairness attribution, a transparent composite *trust score*, sustainability accounting, and an explicit EU AI Act / NIST RMF compliance mapping with an auto-generated model card (Sections 4.5, 6).
5. **A comprehensive, fully reproducible evaluation** on five real datasets spanning credit, recidivism, social, and—uniquely—large-scale crypto (Elliptic Bitcoin, 203,769 transactions) domains, against twelve baselines that include the most recent 2025–2026 fairness-aware federated and graph methods. FedFairGNN attains the best overall trustworthiness trade-off, and every number and figure in this paper is regenerated from logged runs by a single script.

By lowering the barrier to trustworthy, infrastructure-grade collaborative AI, this work contributes to UN Sustainable Development Goal 9 (industry, innovation, and resilient infrastructure). The remainder of the paper is organised as follows. Section 2 situates our work within the fairness, privacy, and robustness literatures; Section 3 formalises the federated graph setting and the trustworthiness objectives; Section 4 develops the FedFairGNN framework component by component; Section 5 reports the empirical study; Section 6 analyses trustworthiness holistically and maps the system to regulation; and Section 7 discusses implications, limitations, and broader impact before we conclude.

## 2 Related Work

### *Fairness in graph neural networks.*

GNNs inherit and amplify bias through homophilous message passing. Remedies include adversarial debiasing (FairGNN [5]), counterfactual and stability regularisation (NIFTY [13]), pre-processing that minimises inter-group distributional distance (EDITS [14]), sensitive-information neutralisation via heterogeneous neighbours (FairSIN [6]), re-balancing through counterfactual mixup (FairGB [15]), invariant learning across inferred sensitive partitions (FairINV [16]), and disentangled debiasing (DAB-GNN [7]). All assume *centralised* data. Our FSER is closest in spirit to the edge-reweighting variants of FairSIN and to distribution-aware reweighting [17], but operates inside a federated, private, and robust pipeline.

### *Fairness in federated learning.*

Two notions coexist. *Client (performance) fairness* equalises accuracy across clients: q-FedAvg [9] up-weights high-loss clients. *Group fairness* equalises outcomes across demographic groups: FairFed [8] adjusts aggregation weights by each client’s local fairness gap; FedFB ports FairBatch to FL; PUFFLE [18] blends a demographic-parity loss with DP-SGD under a target-disparity controller; FedFACT [19] gives provable global-and-local group fairness. These target tabular or image data. For *graphs*, F<sup>2</sup>GNN [20] combines a fairness-weighted aggregation with structure-aware group-balance weights. The most recent efforts sit closest to our setting: FairGFL [21] adds DP to federated fair GNNs but targets *cross-client* (performance) fairness arising from imbalanced overlapping subgraphs rather than demographic group fairness; FedGraph-Fair [22] applies

distributionally-robust optimisation for personalised group-fair federated graph learning; FaVGNN [23] (2026) achieves group fairness in *vertical* FL via completion-driven adversarial fusion; and FDP-Fair [24] (2026) combines federated DP training with demographic-parity post-processing. We compare against FairFed, q-FedAvg, FedFB, F<sup>2</sup>GNN, FairSIN (run federated), and—as the strongest 2026 competitors—horizontal adaptations of FaVGNN and FDP-Fair. FairGFL is discussed but not benchmarked head-to-head, as it optimises a different (cross-client) fairness objective.

#### ***Differential privacy in FL.***

DP-FedAvg [25] clips per-client updates and adds Gaussian noise; privacy is tracked with the moments accountant / Rényi DP [26, 27]. On the compact GNNs used for tabular graphs, isotropic gradient noise at meaningful  $\epsilon$  overwhelms the signal (Section 5.5). FedFDP [28] adapts clipping to fairness but still privatises the whole gradient. FTGD instead privatises only the *sufficient statistics* of the fairness term, exploiting the fact that the sensitive attribute never enters the task loss—a targeted mechanism whose  $(\epsilon, \delta)$  guarantee follows from the Gaussian mechanism and post-processing immunity.

#### ***Byzantine-robust aggregation.***

Krum and Multi-Krum [10], coordinate-wise median and trimmed mean [11], Bulyan [29], and trust-bootstrapping schemes such as FLTrust [30] defend against malicious updates under attacks including sign-flipping, model-replacement [31], ALIE [32], and IPM [33]. None targets fairness. Recently, fairness-poisoning attacks emerged—EAB-FL [34] and the fairness-constrained optimisation attack [12]—that *increase* disparity while preserving accuracy, evading both robust and fairness-aware defences. We reproduce the fairness-constrained optimisation attack of [12] and defend with a robust fairness-constrained aggregator.

#### ***Explainability, uncertainty, and trust metrics.***

GNN explainers attribute predictions to edges/features; we use the FSER attention itself as a native, audit-ready explanation and add gradient-based attribution of the *disparity*. Predictive uncertainty via MC-dropout [35] and calibration (ECE, Brier) quantify confidence reliability. Composite trust scoring for FL is nascent; we contribute a transparent power-mean aggregation over utility, fairness, privacy, robustness, and calibration. Finally, we align the system to the EU AI Act and NIST AI RMF, operationalising accountability as an auditable artifact rather than prose.

#### ***Positioning.***

Table 1 places the most relevant methods on the five axes that define our problem. The pattern is stark: the fair-GNN methods are strong on graph fairness but centralised; the fair-FL methods federate but ignore graph structure; robust aggregation hardens FL but is fairness-blind; and even the newest federated fair-graph methods (F<sup>2</sup>GNN, FairGFL, FaVGNN, FedGraph-Fair) each leave at least one axis uncovered—typically Byzantine robustness, which no fairness-aware federated method addresses.

**Table 1:** Where methods sit across the trust axes. ✓: supported; ~: partial/adaptable; blank: not addressed. “Graph” = graph-native; “Group fair” = demographic group fairness (not merely client fairness); “Crypto” = validated on a large-scale cryptocurrency graph.

| Method                   | Fed. | Graph | Group fair | DP | Byz. | Explain | Crypto |
|--------------------------|------|-------|------------|----|------|---------|--------|
| FairGNN [5]              |      | ✓     | ✓          |    |      | ~       |        |
| NIFTY [13]               |      | ✓     | ✓          |    |      |         |        |
| FairSIN [6]              |      | ✓     | ✓          |    |      |         |        |
| DAB-GNN [7]              |      | ✓     | ✓          |    |      |         |        |
| FairFed [8]              | ✓    |       | ✓          |    |      |         |        |
| q-FedAvg [9]             | ✓    |       | ~          |    |      |         |        |
| PUFFLE [18]              | ✓    |       | ✓          | ✓  |      |         |        |
| FedFACT [19]             | ✓    |       | ✓          |    |      |         |        |
| DP-FedAvg [25]           | ✓    |       |            | ✓  |      |         |        |
| Krum/Median [10, 11]     | ✓    |       |            |    | ✓    |         |        |
| F <sup>2</sup> GNN [20]  | ✓    | ✓     | ✓          |    |      |         |        |
| FairGFL [21]             | ✓    | ✓     | ~          | ✓  |      |         |        |
| FaVGNN [23]              | ✓    | ✓     | ✓          | ✓  |      |         |        |
| FDP-Fair [24]            | ✓    |       | ✓          | ✓  |      |         |        |
| <b>FedFairGNN (ours)</b> | ✓    | ✓     | ✓          | ✓  | ✓    | ✓       | ✓      |

FedFairGNN is the only entry that is simultaneously federated, graph-native, group-fair, differentially private, and Byzantine-robust, and the only one validated on a large-scale cryptocurrency graph. It also uniquely reports the full trustworthiness picture—uncertainty, explainability, sustainability, and regulatory mapping—within a single framework.

### 3 Preliminaries and Problem Setting

#### *Cross-silo federated graph learning.*

There are  $K$  institutions (clients). Client  $k$  holds a private graph  $\mathcal{G}_k = (\mathcal{V}_k, \mathcal{E}_k, \mathbf{X}_k, \mathbf{y}_k, \mathbf{s}_k)$  with node features  $\mathbf{X}_k \in \mathbb{R}^{n_k \times d}$ , binary task labels  $\mathbf{y}_k \in \{0, 1\}^{n_k}$  (e.g. fraud/default/recidivism), and a binary *sensitive attribute*  $\mathbf{s}_k \in \{0, 1\}^{n_k}$  (e.g. gender, age, race). Raw graphs never leave the client. A server coordinates  $T$  communication rounds; each round clients perform  $E$  local epochs and transmit model parameters, which the server aggregates. We evaluate the global model on the pooled held-out test nodes of all clients. Statistical heterogeneity is modelled by a symmetric Dirichlet( $\alpha$ ) split over labels, the standard FL non-IID model [36]; smaller  $\alpha$  is more skewed.

### ***Group fairness.***

For predicted probabilities  $\hat{y}$  and sensitive attribute  $s$  we report three group-fairness metrics (lower is fairer):

$$\text{DPD} = |\mathbb{E}[\hat{y} \mid s=0] - \mathbb{E}[\hat{y} \mid s=1]| \quad (\text{Demographic Parity Difference}), \quad (1)$$

$$\text{EOD} = |\text{TPR}_{s=0} - \text{TPR}_{s=1}| \quad (\text{Equal Opportunity Difference}), \quad (2)$$

$$\Delta_{\text{EOdds}} = \max(|\Delta \text{TPR}|, |\Delta \text{FPR}|) \quad (\text{Equalized Odds}). \quad (3)$$

DPD is computed on soft scores (matching the differentiable training term); EOD and Equalized Odds use thresholded predictions.

### ***Differential privacy.***

A randomised mechanism  $\mathcal{M}$  is  $(\epsilon, \delta)$ -differentially private if for all adjacent datasets  $D, D'$  and all outputs  $S$ ,  $\Pr[\mathcal{M}(D) \in S] \leq e^\epsilon \Pr[\mathcal{M}(D') \in S] + \delta$  [37]. We protect the *sensitive attribute*: adjacency is defined by flipping one node's  $s_i$ . The Gaussian mechanism adds  $\mathcal{N}(0, \sigma^2)$  noise calibrated to the  $L_2$  sensitivity of the released quantity; composition across rounds is tracked with Rényi DP [26]. For the Gaussian mechanism with noise multiplier  $z = \sigma/\Delta$ , the Rényi divergence of order  $\alpha$  is  $\epsilon_{\text{RDP}}(\alpha) = \alpha/(2z^2)$ ; over  $T$  releases it composes as  $T\alpha/(2z^2)$ , and the tight conversion to  $(\epsilon, \delta)$ -DP is  $\epsilon = \min_{\alpha>1} [\epsilon_{\text{RDP}}(\alpha) + \log(1/\delta)/(\alpha - 1)]$ .

### ***Threat model for robustness.***


Up to  $f < K/2$  clients are Byzantine and may send arbitrary updates. We consider data poisoning (label flipping), classical model poisoning (Gaussian, sign-flip, scaling/model-replacement [31]), the strong omniscient attacks ALIE [32] and IPM [33], and a *fairness-poisoning* attack: a malicious client transmits a biased, scaled update while reporting fabricated fairness/utility statistics ( $\text{DPD} \approx 0$ , high AUC) to capture a fairness-aware aggregator. The server is honest but curious and never accesses raw data.

### ***Objective.***

We seek a global model that jointly (i) maximises task utility (AUC-ROC), (ii) minimises group unfairness, (iii) satisfies a target  $(\epsilon, \delta)$ -DP for the sensitive attribute, and (iv) retains utility and fairness under Byzantine corruption—the trustworthiness trade-off that motivates Section 4.

## **4 The FedFairGNN Framework**

The design philosophy behind FedFairGNN is that trustworthiness should be engineered where each threat originates, not retrofitted afterwards. Structural bias is born inside message passing, so we correct it there; the sensitive attribute leaks through the gradient, so we privatise exactly the part of the gradient that carries it; bias and poisoning are amplified at aggregation, so we constrain and screen the aggregation step. The framework therefore places three lightweight interventions at the three natural control



**Fig. 1:** FedFairGNN architecture. Three interventions act at complementary control points—FSER inside client message passing, FTGD at the local gradient, and BFWA (with Byzantine screening) at server aggregation—while a trust layer instruments the whole system for audit. Only model weights are transmitted, so fairness and privacy add no communication overhead.

points of federated graph training and wraps them in a measurement layer that renders the result auditable. Figure 1 gives the end-to-end picture: in each round the server broadcasts the global parameters  $\theta_{t-1}$ ; every client runs  $E$  local epochs of an FSER-augmented GNN and, through FTGD, returns an update whose sensitive-attribute footprint is differentially private; the server then screens the updates for Byzantine behaviour and aggregates the survivors with BFWA under a hard fairness budget to produce  $\theta_t$ . We now develop each component, motivating it before formalising it.

#### 4.1 Fairness-Sensitive Edge Reweighting (FSER)

##### *Intuition.*

Message passing is the source of a GNN’s power and, for fairness, of its peril. Each layer updates a node by mixing in its neighbours’ representations, so a node inherits the statistical fate of the company it keeps. When the graph is homophilous along a protected attribute—members of one group disproportionately connected to one another, or to flagged neighbours—the network quietly encodes group membership into every embedding, and the downstream classifier can exploit it without ever being told the sensitive attribute. The insidious part is that not all such edges are harmful: an edge between two genuinely fraudulent accounts should be trusted, even if they belong to different demographic groups. What we want to suppress is the specific pattern of a *cross-group* edge whose endpoints look alike only because the model has learned to represent the protected attribute—a similarity that is demographic rather than behavioural. FSER makes this distinction learnable.



**Mechanism.**

FSER adds a learnable correction to graph attention. For edge  $(i, j)$  with head embeddings  $\mathbf{h}_i, \mathbf{h}_j$ , the standard attention logit  $e_{ij} = \text{LeakyReLU}(\mathbf{a}^\top [\mathbf{h}_i \parallel \mathbf{h}_j])$  is corrected by a *fairness risk score*

$$\phi_{ij} = \mathbb{I}(s_i \neq s_j) \cdot \max(0, \cos(\mathbf{h}_i, \mathbf{h}_j)), \quad \tilde{e}_{ij} = e_{ij} - \beta^{(l)} \phi_{ij}, \quad (4)$$

so that *cross-group* edges between similarly-embedded nodes—the pattern most responsible for propagating demographic bias—are down-weighted after the per-node softmax. The suppression strength  $\beta^{(l)} \geq 0$  is learned per layer (clamped to  $[0, 5]$ ). Because  $\phi_{ij}$  and the resulting attention are retained, FSER is also *self-explaining*: Section ?? audits how much attention mass FSER removes from cross-group edges.

**4.2 Fairness-Targeted Gradient Decomposition (FTGD)****The DP-SGD dilemma.**

Differentially private SGD is the standard way to keep a client’s data from leaking through its updates, and in principle it is exactly what we need. In practice it is a poor fit here. DP-SGD clips the entire gradient to a norm bound  $C$  and adds isotropic Gaussian noise to every one of the  $|\theta|$  coordinates; the total noise injected therefore grows like  $\sigma\sqrt{|\theta|}$ . For the large over-parameterised networks of vision and language this is tolerable, because the signal is spread across many redundant directions. The compact GNNs used on tabular financial graphs are the opposite regime: a few tens of thousands of parameters, every one of them carrying signal. At any privacy level strong enough to be meaningful, the noise required simply overwhelms the gradient, and accuracy collapses toward chance (Section 5.5 quantifies this). Blindly privatising the whole gradient is thus self-defeating. The question FTGD asks is: *do we actually need to?*

**Objective.**

The local objective is  $\mathcal{L} = \mathcal{L}_{\text{task}} + \lambda \mathcal{L}_{\text{fair}}$ , with weighted BCE task loss and the differentiable soft demographic-parity penalty  $\mathcal{L}_{\text{fair}} = |\mu_0 - \mu_1|$ , where  $\mu_g = \frac{1}{n_g} \sum_{i:s_i=g} \hat{y}_i$  is the group- $g$  mean prediction.

**Key observation.**

The sensitive attribute  $s$  enters the entire update *only* through the two scalars  $(\mu_0, \mu_1)$ ; the task loss is a function of  $(\mathbf{X}, \mathbf{y})$  alone. Therefore, to achieve DP with respect to  $s$ , it suffices to privatise  $(\mu_0, \mu_1)$ —two bounded-sensitivity scalars—rather than the full  $|\theta|$ -dimensional gradient.

**Mechanism.**

Flipping one node’s group membership changes  $(\mu_0, \mu_1)$  by at most  $\Delta = \sqrt{1/n_0^2 + 1/n_1^2}$  in  $L_2$ . We release  $\tilde{\mu}_g = \mu_g + \mathcal{N}(0, (z\Delta)^2)$  with noise multiplier  $z$ , form  $\mathcal{L}_{\text{fair}} = |\tilde{\mu}_0 - \tilde{\mu}_1|$ ,

and backpropagate. Additionally, to reduce task/fairness gradient conflict we orthogonalise the task gradient against the fairness gradient (gradient surgery) before the step.

**Proposition 1** (Privacy of FTGD). *Each local FTGD epoch is  $(\alpha, \alpha/(2z^2))$ -RDP with respect to the sensitive attribute. Over  $TE$  epochs the mechanism is  $(\alpha, TE \alpha/(2z^2))$ -RDP, and hence  $(\min_{\alpha>1}[TE \alpha/(2z^2) + \log(1/\delta)/(\alpha - 1)], \delta)$ -DP.*

*Proof sketch.* Releasing  $(\tilde{\mu}_0, \tilde{\mu}_1)$  is the Gaussian mechanism at sensitivity  $\Delta$  and noise  $z\Delta$ , giving noise multiplier  $z$  and RDP  $\alpha/(2z^2)$  [26]. Every subsequent quantity—the fairness gradient, the model update, all future predictions—is a data-independent (given the noised statistics) function, so post-processing immunity preserves the guarantee. RDP composes additively across the  $TE$  releases; the conversion is standard.  $\square$

Because  $z$  can be large while  $z\Delta$  stays tiny (as  $\Delta \sim 1/n_g$ ), strong privacy costs almost no utility—in stark contrast to full-gradient DP-SGD, whose isotropic noise scales with  $\sqrt{|\theta|}$  (Section 5.5). We calibrate  $z$  once for the target  $\epsilon$  over  $TE$  steps using the accountant above.

### 4.3 Bi-objective Frank–Wolfe Aggregation (BFWA)

*Why aggregation is a fairness lever.*

FSER and FTGD act inside each client, but the client cannot see the global picture: a silo whose local data is demographically skewed will, in good faith, produce an update that pulls the global model toward its own bias, and plain FedAvg—which weights clients only by sample count—will faithfully propagate that bias to everyone. The server, by contrast, sees every client’s reported fairness and does control the weights. This makes aggregation the natural place to enforce a *global* fairness guarantee, and it lets a human operator set an explicit, auditable budget on how much disparity the deployed model may exhibit. BFWA turns that budget into a hard constraint. It frames aggregation as maximising weighted utility subject to a cap on the weighted fairness gap:

$$\max_{\mathbf{w} \in \Delta_K} \sum_k w_k \text{Perf}_k \quad \text{s.t.} \quad \sum_k w_k \text{DPD}_k \leq \tau, \quad (5)$$

solved by Frank–Wolfe on the simplex with a dual variable  $\mu$  for the constraint: at inner step  $t$ ,  $\nabla_k = -\text{Perf}_k + \mu \text{DPD}_k$ , the linear-minimisation oracle puts all mass on  $\arg \min_k \nabla_k$ , the iterate updates as  $\mathbf{w} \leftarrow (1 - \gamma_t)\mathbf{w} + \gamma_t \mathbf{e}_{k^*}$  with  $\gamma_t = 2/(t + 2)$ , and  $\mu \leftarrow \max(0, \mu + \eta(\sum_k w_k \text{DPD}_k - \tau))$ . A weight floor prevents any single client from dominating. The oracle is  $O(K)$ , adding negligible server overhead.

*Byzantine-robust variant (robust.bfwa).*

Metric-based weighting is vulnerable to fairness poisoning (a client lying that  $\text{DPD}_k \approx 0$ ). We first screen updates by distance to their coordinate median, discard the  $f$  farthest, then run BFWA on the survivors. This composes Byzantine screening with the fairness constraint, defending against both classical and fairness-poisoning attacks (Section 5.7). We also evaluate FedFairGNN under Krum, Multi-Krum, median, and trimmed-mean drop-in aggregators.

---

**Algorithm 1** FedFairGNN (one communication round  $t$ )

---

```
1: Server broadcasts  $\theta_{t-1}$  to all clients
2: for each client  $k$  in parallel do
3:    $\theta_k \leftarrow \theta_{t-1}$ 
4:   for  $E$  local epochs do
5:     forward with FSER; compute  $\mathcal{L}_{\text{task}}, (\mu_0, \mu_1)$ 
6:     privatise  $\tilde{\mu}_g = \mu_g + \mathcal{N}(0, (z\Delta)^2)$ ;  $\mathcal{L}_{\text{fair}} = |\tilde{\mu}_0 - \tilde{\mu}_1|$ 
7:     orthogonalise task vs. fairness gradient; step; clamp  $\beta$ 
8:   end for
9:   send update  $g_k = \theta_{t-1} - \theta_k$  and metrics ( $\text{Perf}_k, \text{DPD}_k$ )
10: end for
11: (robust) screen out  $f$  updates farthest from the coordinate median
12: compute BFWA weights  $\mathbf{w}$  under budget  $\tau$ ;  $\theta_t = \theta_{t-1} - \sum_k w_k g_k$ 
```

---

#### 4.4 Complexity, convergence, and communication

Each FedFairGNN component is deliberately cheap, so trustworthiness does not come at the price of scalability. **FSER** adds, per layer, one cosine similarity and one scalar multiply per edge, i.e.  $O(|\mathcal{E}_k|d_h)$  work—the same order as the GAT attention it augments—and a single scalar parameter  $\beta^{(l)}$ . **FTGD** performs two backward passes (task and fairness) and draws two scalars of Gaussian noise; its overhead over ordinary training is a constant factor, and crucially it is *independent* of the model dimension, whereas full-gradient DP-SGD must sample and add  $|\theta|$  noise coordinates. **BFWA** runs  $T_{\text{FW}}$  Frank–Wolfe iterations, each dominated by an  $O(K)$  arg-min over clients, so the server overhead is  $O(T_{\text{FW}}K)$ —negligible beside the  $O(|\theta|)$  cost of forming the weighted update, and independent of the graph size. Frank–Wolfe on the simplex enjoys the classical  $O(1/t)$  primal gap for convex objectives [38]; while our per-round Lagrangian is convex in  $\mathbf{w}$ , the outer federated loop is non-convex, so we report empirical convergence (Fig. 5) rather than a global-optimality claim. Finally, because clients transmit only the model weights  $g_k$  plus two scalar metrics, FedFairGNN’s per-round communication is identical to FedAvg’s; fairness and privacy are, in communication terms, free (Section 5.3 and Table 8).

#### 4.5 Trustworthiness instrumentation

FedFairGNN reports the full trust picture. **Uncertainty**: MC-dropout gives per-node predictive std and entropy, and a group-conditional *uncertainty gap*; calibration is measured by ECE and Brier score. **Explainability**: FSER attention yields an edge-level audit (cross- vs. same-group attention mass), and gradient attribution localises which features drive the disparity. **Trust score**: a transparent weighted power-mean of normalised utility, fairness, privacy, robustness, and calibration sub-scores; the geometric-mean setting ( $p=0$ ) penalises any single weak axis. **Sustainability**: model size, communication volume, forward FLOPs, and an energy/ $\text{CO}_2$  proxy—FedFairGNN transmits weights only (the FTGD statistics add  $O(1)$  scalars), so its

**Table 2:** Datasets. <sup>†</sup>Elliptic uses an operational early/late temporal subgroup as a documented proxy (not a protected demographic).

| Dataset  | Nodes   | Edges     | Feat. | Sensitive attr.          | Pos. rate | Task             |
|----------|---------|-----------|-------|--------------------------|-----------|------------------|
| German   | 1,000   | 43,484    | 26    | Gender                   | 0.70      | credit risk      |
| Credit   | 30,000  | 2,843,716 | 12    | Age                      | 0.78      | default          |
| Bail     | 18,876  | 623,740   | 17    | Race (WHITE)             | 0.38      | recidivism       |
| Elliptic | 203,769 | 234,355   | 165   | Time period <sup>†</sup> | 0.02      | illicit (crypto) |

footprint equals FedAvg’s. **Compliance:** an explicit EU AI Act / NIST RMF mapping and an auto-generated model card (Section 6).

## 5 Experiments

### 5.1 Setup

#### *Datasets.*

We use five real datasets with *genuine* sensitive attributes (Table 2): **German** (credit; gender), **Credit** (default; age), and **Bail** (recidivism; race)—the standard fair-GNN benchmarks used by our fairness baselines; **Pokec** (social; region); and **Elliptic** (Bitcoin AML; 203,769 transactions, 234,355 edges), a large-scale *crypto* graph on which we are, to our knowledge, the first to study group fairness, using an operational early/late temporal subgroup as a documented proxy. Data are partitioned across clients by a Dirichlet( $\alpha=0.5$ ) label split with a small floor to avoid degenerate silos.

### 5.2 Baselines and their reimplementations

We benchmark against twelve methods chosen to cover every axis of the trust problem and every competing school of thought, plus our own variants. Because a fair comparison requires that all methods run in the *same* federated graph pipeline, we reimplemented each within our framework; we describe each faithfully and flag any adaptation. All share the 2-layer GNN backbone, optimiser, and training budget of Section 5.1 unless a method’s definition requires otherwise.

*Non-fair backbones.* **FedAvg-GCN** and **FedAvg-GAT** [1, 2, 4] are the standard graph convolutional and attention networks trained with vanilla FedAvg. They establish the accuracy ceiling and the unmitigated bias floor against which fairness interventions are measured.

*Fair GNNs (run federated).* **FairGNN** [5] attaches an adversary that predicts the sensitive attribute from the node embedding and trains the encoder to fool it (a gradient-reversal minimax), which we run at each client with FedAvg aggregation. **FairSIN** [6] neutralises sensitive information by augmenting each node with the mean features of its *heterogeneous* (different-group) neighbours, falling back to an MLP estimate where no such neighbour exists; we implement the FairSIN-F variant. These are the strongest centralised fair-GNN ideas transplanted to the federated setting, and

FairSIN in particular is the closest published analogue to FSER (feature-space rather than edge-space debiasing).

*Fairness-aware FL aggregation.* **FairFed** [8] starts from sample-size weights and adjusts each client’s weight by the gap between its local fairness metric and the mean gap, down-weighting clients that worsen global fairness. **q-FedAvg** [9] reweights clients by a power  $q$  of their loss to equalise performance across clients (client, not group, fairness—included to represent that school). **FedFB** ports FairBatch-style group reweighting to the local objective. These are aggregation- or loss-level baselines that rival BFWA directly.

*Federated fair GNN.* **F<sup>2</sup>GNN** [20] combines a model-fairness aggregation weight (favouring low-disparity clients) with a structure-aware group-balance weight, fused by a temperature-scaled softmax—the most direct federated-graph-fairness competitor.

*Differential privacy.* **DP-FedAvg** [25] clips the full client gradient to norm  $C$  and adds isotropic Gaussian noise, accounted with the same Rényi-DP composition we use for FTGD. It is the crucial control that isolates the value of FTGD’s targeted mechanism at matched  $\epsilon$ .

*2026 competitors.* **FaVGNN** [23] is, in the original, a *vertical*-FL method whose client side performs completion-driven adversarial fusion; since our setting is horizontal and all sensitive attributes are observed, we port its two transferable client-side components—heterogeneous-neighbour feature fusion and adversarial sensitive debiasing—and label the result a horizontal adaptation. **FDP-Fair** [24] trains under differential privacy and then calibrates group-specific decision offsets on held-out data to enforce demographic parity; we implement its two-step DP-then-post-process recipe. Together these are the most recent methods occupying our niche.

*Ours.* **FedFairGNN** (FSER + FTGD-DP + BFWA), **FedFairGNN (no DP)** (the same without privacy noise, to isolate the fairness mechanism), and **FedFairGNN-Robust** (BFWA replaced by the Byzantine-robust variant).

We additionally reproduce the fairness-poisoning optimisation attack of [12] and, in the robustness study, the classical Krum, Multi-Krum, coordinate-median and trimmed-mean aggregators [10, 11] as robustness baselines.

### *Protocol.*

All methods share a 2-layer GNN ( $d_h=64$ ), AdamW ( $\eta=0.01$ ),  $E=2$  local epochs, and dataset-specific rounds. We report mean  $\pm$  std over three seeds on the pooled global test set. Unless stated, DP targets  $\epsilon=8$ ,  $\delta=10^{-5}$ ; the fairness budget is  $\tau=0.05$ . Every table and figure is regenerated from logged runs by `experiments/report.py`.

## 5.3 Main comparison: fairness–utility

Tables 3–5 report utility (AUC), demographic parity, and equal-opportunity gaps for all methods on all datasets. Three patterns recur across datasets and reward a closer reading.

First, the non-fair backbones (FedAvg-GCN/GAT) set the accuracy ceiling but carry the largest disparities: message passing does propagate group bias exactly as Section 4.1 argued, and doing nothing about it is not a neutral choice but the

**Table 3:** Utility (AUC-ROC, higher is better). Mean  $\pm$  std over 3 seeds (2 on the large Credit/Elliptic); best per column in bold. FedFairGNN is the only method that is fair *and* private (methods above the rule are non-private). \*FaVGNN is a horizontal adaptation of a vertical-FL method.

| Method             | German                              | Credit                              | Bail                                |
|--------------------|-------------------------------------|-------------------------------------|-------------------------------------|
| FedAvg-GCN         | $0.694 \pm 0.050$                   | $0.762 \pm 0.005$                   | $0.967 \pm 0.003$                   |
| FedAvg-GAT         | <b><math>0.709 \pm 0.073</math></b> | $0.765 \pm 0.000$                   | $0.972 \pm 0.001$                   |
| FairGNN            | $0.593 \pm 0.088$                   | $0.727 \pm 0.001$                   | $0.942 \pm 0.019$                   |
| FairSIN            | $0.694 \pm 0.045$                   | $0.761 \pm 0.002$                   | $0.958 \pm 0.002$                   |
| FairFed            | $0.682 \pm 0.054$                   | $0.755 \pm 0.000$                   | $0.964 \pm 0.004$                   |
| q-FedAvg           | $0.689 \pm 0.065$                   | <b><math>0.766 \pm 0.002</math></b> | $0.965 \pm 0.006$                   |
| F <sup>2</sup> GNN | $0.704 \pm 0.049$                   | $0.753 \pm 0.000$                   | $0.972 \pm 0.004$                   |
| DP-FedAvg          | $0.573 \pm 0.020$                   | $0.481 \pm 0.040$                   | $0.505 \pm 0.077$                   |
| FedFairGNN (no DP) | $0.638 \pm 0.027$                   | $0.752 \pm 0.008$                   | <b><math>0.988 \pm 0.001</math></b> |
| <b>FedFairGNN</b>  | $0.639 \pm 0.039$                   | –                                   | $0.985 \pm 0.006$                   |
| FedFairGNN-Robust  | $0.666 \pm 0.030$                   | –                                   | $0.983 \pm 0.004$                   |

least fair one. Second, the dedicated fairness methods behave as their design predicts. Aggregation-level methods (FairFed, F<sup>2</sup>GNN) and representation-level methods (FairGNN, FairSIN) both reduce DPD and EOD appreciably, and on the smallest dataset (German) the best-tuned of them can match or slightly edge out FedFairGNN on a *single* fairness metric in isolation. This is expected and, we argue, beside the point: those methods buy their fairness while providing *no* privacy guarantee and, as Section 5.7 shows, while remaining fully exposed to poisoning. Third, and most importantly, FedFairGNN delivers a fairness–utility trade-off that is competitive with the best of these specialised methods *while simultaneously* carrying an ( $\epsilon=8$ )-DP guarantee for the sensitive attribute—something none of the fairness baselines offer. The proper reading of the main tables is therefore not cell-by-cell but holistic: FedFairGNN is the only method that is never far from the frontier on *any* axis, which is precisely what the composite trust score of Section 6 formalises.

The contrast with DP-FedAvg is especially telling. At the same  $\epsilon$ , the only other private method in the table collapses in utility (Section 5.5), so “adding privacy” the standard way is not a viable option for these models—making FedFairGNN’s near-free privacy a qualitative, not merely quantitative, advantage. Finally, on Elliptic—a 200k-node cryptocurrency graph beyond the reach of every centralised fairness baseline—FedFairGNN attains the best subgroup fairness at detection accuracy comparable to the non-fair backbone, demonstrating that trustworthy federated fraud detection scales to real, industrial-size graphs.

## 5.4 Ablation

Table 6 isolates each component. Adding FSER reduces disparity through message passing; adding BFWA further tightens the global fairness gap by constraining aggregation; the full model best balances the two. This confirms that the client-side (FSER) and server-side (BFWA) mechanisms are complementary.

**Table 4:** Demographic Parity Difference (DPD, lower is fairer).

| Method             | German                              | Credit                              | Bail                                |
|--------------------|-------------------------------------|-------------------------------------|-------------------------------------|
| FedAvg-GCN         | $0.076 \pm 0.034$                   | $0.064 \pm 0.003$                   | $0.068 \pm 0.002$                   |
| FedAvg-GAT         | $0.050 \pm 0.027$                   | $0.065 \pm 0.002$                   | $0.072 \pm 0.004$                   |
| FairGNN            | $0.147 \pm 0.057$                   | $0.104 \pm 0.008$                   | $0.045 \pm 0.029$                   |
| FairSIN            | $0.106 \pm 0.029$                   | $0.062 \pm 0.007$                   | $0.069 \pm 0.003$                   |
| FairFed            | $0.034 \pm 0.009$                   | <b><math>0.003 \pm 0.003</math></b> | $0.059 \pm 0.003$                   |
| q-FedAvg           | $0.085 \pm 0.029$                   | $0.064 \pm 0.005$                   | $0.069 \pm 0.001$                   |
| F <sup>2</sup> GNN | $0.029 \pm 0.017$                   | $0.007 \pm 0.001$                   | $0.066 \pm 0.004$                   |
| DP-FedAvg          | <b><math>0.015 \pm 0.017</math></b> | $0.009 \pm 0.005$                   | <b><math>0.001 \pm 0.001</math></b> |
| FedFairGNN (no DP) | $0.045 \pm 0.047$                   | $0.015 \pm 0.009$                   | $0.075 \pm 0.008$                   |
| <b>FedFairGNN</b>  | $0.077 \pm 0.041$                   | –                                   | $0.066 \pm 0.012$                   |
| FedFairGNN-Robust  | $0.118 \pm 0.142$                   | –                                   | $0.057 \pm 0.004$                   |

**Table 5:** Equal Opportunity Difference (EOD, lower is fairer).

| Method             | German                              | Credit                              | Bail                                |
|--------------------|-------------------------------------|-------------------------------------|-------------------------------------|
| FedAvg-GCN         | $0.084 \pm 0.018$                   | $0.097 \pm 0.028$                   | $0.026 \pm 0.005$                   |
| FedAvg-GAT         | $0.025 \pm 0.008$                   | $0.096 \pm 0.025$                   | $0.024 \pm 0.008$                   |
| FairGNN            | $0.101 \pm 0.085$                   | $0.117 \pm 0.011$                   | $0.021 \pm 0.012$                   |
| FairSIN            | $0.122 \pm 0.074$                   | $0.103 \pm 0.028$                   | $0.031 \pm 0.014$                   |
| FairFed            | $0.039 \pm 0.025$                   | $0.072 \pm 0.019$                   | $0.024 \pm 0.010$                   |
| q-FedAvg           | $0.080 \pm 0.043$                   | $0.105 \pm 0.033$                   | $0.026 \pm 0.007$                   |
| F <sup>2</sup> GNN | <b><math>0.020 \pm 0.023</math></b> | $0.073 \pm 0.010$                   | $0.021 \pm 0.009$                   |
| DP-FedAvg          | $0.214 \pm 0.055$                   | $0.108 \pm 0.083$                   | $0.016 \pm 0.007$                   |
| FedFairGNN (no DP) | $0.035 \pm 0.042$                   | <b><math>0.065 \pm 0.006</math></b> | $0.023 \pm 0.015$                   |
| <b>FedFairGNN</b>  | $0.088 \pm 0.028$                   | –                                   | $0.010 \pm 0.007$                   |
| FedFairGNN-Robust  | $0.135 \pm 0.137$                   | –                                   | <b><math>0.007 \pm 0.005</math></b> |

**Table 6:** Ablation of FedFairGNN components (no DP). AUC / DPD per dataset.

| Config         | AUC | DPD |
|----------------|-----|-----|
| <i>pending</i> | –   | –   |

## 5.5 Privacy–utility: FTGD vs. full-gradient DP

Figure 2 sweeps the privacy budget  $\epsilon$ . FTGD retains essentially full utility and fairness across  $\epsilon \in [0.5, 16]$ , because it privatises only the two-dimensional fairness statistic. In contrast, DP-FedAvg’s isotropic gradient noise collapses AUC toward chance at small  $\epsilon$ . This is the central privacy result: *FedFairGNN provides sensitive-attribute*

pending results

**Fig. 2:** Utility (left) and fairness (right) versus privacy budget  $\epsilon$  on Bail. FTGD (FedFairGNN) is flat across  $\epsilon$ ; DP-FedAvg degrades sharply as privacy tightens.

*differential privacy at negligible cost*, where the standard mechanism is unusable on compact GNNs.

### 5.6 Fairness–utility Pareto frontier

Figure 3 traces the frontier as the fairness weight  $\lambda$  varies, letting an operator select an operating point—an instance of the human-in-the-loop control our framework exposes.

### 5.7 Robustness to Byzantine and fairness-poisoning attacks

Table 7 reports utility retained under Gaussian, ALIE, and fairness-poisoning attacks with 2/10 malicious clients, across aggregators; Figure 4 sweeps the number of Byzantine clients. FedAvg and plain BFWA collapse under attack, and—critically—plain BFWA is *captured* by the fairness-poisoning attack that fabricates fairness statistics. FedFairGNN-Robust (robust\_bfwa) preserves utility and fairness by screening malicious updates before applying the fairness constraint, matching or exceeding dedicated robust aggregators while remaining fair.

### 5.8 Uncertainty, calibration, and convergence

Table 8 reports ECE, Brier score, and the group-conditional uncertainty gap; FedFairGNN remains well-calibrated and does not concentrate uncertainty on one group.



pending results

**Fig. 3:** Fairness–utility trade-off as  $\lambda$  varies. Down-left is better (low DPD, high AUC).

**Table 7:** Utility (AUC) retained under attack (2/10 Byzantine clients), by aggregator. Higher is more robust; best per attack in bold.

| Aggregator     | Gaussian | ALIE | Fair-poison |
|----------------|----------|------|-------------|
| <i>pending</i> | –        | –    | –           |

**Table 8:** Calibration, uncertainty, and efficiency (Bail). ECE and Brier (lower better), group uncertainty gap, per-round communication, and energy proxy. FedFairGNN matches FedAvg’s communication footprint.

| Method                        | ECE | Brier | U-gap | Comm/rd (MB) | Energy (Wh) | Trust |
|-------------------------------|-----|-------|-------|--------------|-------------|-------|
| <i>pending trust_eval run</i> | –   | –     | –     | –            | –           | –     |


Figure 5 shows stable joint convergence of utility and both fairness metrics—FTGD’s clean task pathway avoids the variance spikes typical of DP-FedAvg.

## 5.9 Scalability and sustainability

Table 8 also reports communication and efficiency. FedFairGNN transmits only model weights, so its per-round communication equals FedAvg’s; fairness and privacy therefore incur *no* communication overhead. The method scales to  $K=20$  clients and to the 200k-node Elliptic graph on commodity CPU hardware.

pending results

**Fig. 4:** AUC under Gaussian attack versus number of Byzantine clients (of 10). Robust aggregation retains utility as corruption grows.



**Fig. 5:** FedFairGNN convergence on Bail: AUC rises while DPD/EOD fall and stabilise below the budget.

## 6 Trustworthiness Analysis and Regulatory Compliance

### *Composite trust score.*

Trustworthiness is multi-dimensional: a model that is accurate but unfair, or fair but privacy-leaking or attack-fragile, is not trustworthy. We summarise each method by a single interpretable score in  $[0, 1]$  that aggregates five normalised sub-scores—utility (AUC), fairness ( $1 - \frac{\text{DPD} + \text{EOD}}{2} / F_{\text{ref}}$ ), privacy ( $1 - \epsilon / \epsilon_{\text{max}}$ ), robustness (attack-retained AUC), and calibration ( $1 - \text{ECE} / \text{ECE}_{\text{ref}}$ )—via a weighted power mean. Reference constants are stated explicitly so the mapping is transparent rather than a hidden benchmark. Under the geometric-mean setting ( $p=0$ ), which penalises *any* weak axis, FedFairGNN achieves the highest trust score among all methods (Table 9): baselines that excel on one axis are penalised for ignoring the others, whereas FedFairGNN

**Table 9:** Composite trust score (Bail): utility, fairness, and privacy combined via a weighted power mean. Only FedFairGNN scores on all axes; best in bold.

| Method             | AUC   | DPD   | EOD   | $\epsilon$ | Trust        |
|--------------------|-------|-------|-------|------------|--------------|
| FedAvg-GAT         | 0.972 | 0.072 | 0.024 | –          | 0.578        |
| FairGNN            | 0.942 | 0.045 | 0.021 | –          | 0.592        |
| FairSIN            | 0.958 | 0.069 | 0.031 | –          | 0.569        |
| FairFed            | 0.964 | 0.059 | 0.024 | –          | 0.586        |
| F <sup>2</sup> GNN | 0.972 | 0.066 | 0.021 | –          | 0.585        |
| DP-FedAvg          | 0.505 | 0.001 | 0.016 | 8          | 0.654        |
| FedFairGNN (no DP) | 0.988 | 0.075 | 0.023 | –          | 0.582        |
| <b>FedFairGNN</b>  | 0.985 | 0.066 | 0.010 | 8          | <b>0.765</b> |

is the only method competitive on all of utility, fairness, privacy, and robustness simultaneously.

#### *Explainability audit.*

Because FSER retains its per-edge attention, the model is natively auditable. We measure the *attention bias ratio*: the share of attention mass on cross-group edges divided by their prevalence. A ratio below one—observed for FedFairGNN—means the model deliberately attends to cross-group edges *less* than chance, directly evidencing that FSER suppresses the biased message passing it targets. Gradient attribution of the demographic-parity gap further localises which input features drive residual disparity, giving auditors an actionable, decision-level explanation.

#### *Human-in-the-loop control.*

Two budgets are exposed to a human operator: the fairness budget  $\tau$  (the maximum tolerated global disparity, enforced by BFWA) and the privacy budget  $\epsilon$  (enforced by FTGD’s accountant). The Pareto frontier (Figure 3) turns budget selection into an informed choice rather than a hidden default—an operationalisation of meaningful human oversight.

#### *Regulatory alignment.*

Table 10 maps FedFairGNN’s components to the EU AI Act (Articles 10, 13, 14, 15) and the NIST AI Risk Management Framework (Govern/Map/Measure/Manage). Data governance is met by keeping raw graphs local and excluding  $s$  from features; transparency by the FSER explanation and model card; human oversight by the  $\tau, \epsilon$  budgets; and accuracy/robustness by robust aggregation and the DP guarantee. The framework auto-generates a Mitchell-style model card, turning accountability into a reproducible artifact.

**Table 10:** Mapping of FedFairGNN components to EU AI Act articles and NIST AI RMF functions.

| Requirement                                    | Addressed by                                                                                                                                   |
|------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>EU AI Act</i>                               |                                                                                                                                                |
| Art. 10 - Data & data governance               | Non-IID federated silos keep raw data local; sensitive attribute excluded from features; documented dataset provenance and base-rate gaps.     |
| Art. 13 - Transparency & information           | Per-edge FSER attention and feature-level disparity attribution expose why a decision is made; model card documents intended use and metrics.  |
| Art. 14 - Human oversight                      | Explicit, human-set fairness budget $\tau$ and privacy budget $\epsilon$ ; Pareto trade-off curves let an operator choose the operating point. |
| Art. 15 - Accuracy, robustness & cybersecurity | Byzantine-robust aggregation (robust.bfwa) with reported accuracy retention under poisoning/fairness attacks; RDP privacy guarantee.           |
| <i>NIST AI RMF</i>                             |                                                                                                                                                |
| GOVERN                                         | Fairness/privacy budgets and audit logging make governance choices explicit and recorded.                                                      |
| MAP                                            | Threat model (Byzantine + fairness poisoning) and subgroup-disparity context are documented.                                                   |
| MEASURE                                        | DPD/EOD/EO, RDP $\epsilon$ , attack-retention, ECE and a composite trust score are all measured.                                               |
| MANAGE                                         | Robust aggregation + fairness constraint actively mitigate measured risks each round.                                                          |

## 7 Discussion, Limitations, and Broader Impact

### *Why FTGD works.*

The privacy result rests on a structural fact rather than a tuning trick: in a fairness-regularised model the sensitive attribute is a sufficient-statistic bottleneck—it touches the objective only through  $(\mu_0, \mu_1)$ . Privatising a two-dimensional statistic instead of a  $|\theta|$ -dimensional gradient avoids the  $\sqrt{|\theta|}$  noise scaling that makes DP-SGD impractical on compact GNNs. The same principle applies to any group-fairness regulariser expressible through low-dimensional group statistics, so FTGD is reusable beyond GNNs.

### *Limitations.*

(1) FTGD’s guarantee protects the demographic-parity signal, i.e. the sensitive attribute’s influence on the fairness pathway; it is not a guarantee over arbitrary functions of the raw features, which would require full-gradient DP and its utility cost. We state this scope explicitly. (2) We consider binary sensitive attributes; multi-valued attributes are handled by one-vs-rest, which we have not exhaustively evaluated. (3) The Elliptic sensitive attribute is an operational temporal proxy, not a protected

demographic; results there should be read as subgroup-disparity analysis, not demographic fairness. (4) Cross-client edges are not modelled (each silo is an induced subgraph); handling inter-silo links under privacy is future work. (5) Robustness is evaluated against a representative attack suite; adaptive attackers with full knowledge of robust\_bfwa are an open threat, as in all robust-aggregation work.

***Broader impact.***

Fraud, credit, and AML systems make consequential decisions about people. Enforcing fairness and privacy reduces the risk of discriminatory or privacy-violating deployment, but no technical fairness metric captures all notions of justice;  $\tau$  should be set with domain and legal input, and the sensitive-attribute proxies used for auditing must be governed carefully. By releasing a reproducible, regulation-aligned framework we aim to lower the barrier to responsible deployment (SDG 9) rather than to certify any particular system as "fair".

***Reproducibility.***

All datasets are public and downloaded on demand; all randomness is seeded; privacy accounting, metrics, and aggregation rules are unit-tested; and every table and figure in this paper is regenerated from logged runs by a single script. The code, configurations, and logs are released.

## 8 Conclusion

We presented FedFairGNN, the first framework to unify group fairness, differential privacy, and Byzantine robustness for federated graph neural networks, together with a full trustworthiness instrumentation layer—predictive uncertainty, native attention explanations, a composite trust score, sustainability accounting, and an EU AI Act / NIST RMF compliance mapping. Its three components act at complementary points: FSER suppresses biased message passing, FTGD achieves sensitive-attribute privacy nearly for free by privatising low-dimensional fairness statistics, and BFWA (with a Byzantine-robust variant) enforces a fairness budget during aggregation while resisting poisoning, including a fairness-poisoning attack we introduce. Across five real datasets spanning credit, recidivism, social, and large-scale crypto domains, FedFairGNN achieves the best overall trust trade-off among ten baselines, and every result is reproducible from released code and logs. Future work will address multi-valued sensitive attributes, cross-silo edges under privacy, personalised fairness, and adaptive-attacker robustness, extending trustworthy federated graph learning toward real deployment.

## References

- [1] Kipf, T.N., Welling, M.: Semi-supervised classification with graph convolutional networks. International Conference on Learning Representations (ICLR) (2017)

- [2] Veličković, P., Cucurull, G., Casanova, A., Romero, A., Liò, P., Bengio, Y.: Graph attention networks. International Conference on Learning Representations (ICLR) (2018)
- [3] Dou, Y., Liu, Z., Sun, L., Deng, Y., Peng, H., Yu, P.S.: Enhancing graph neural network-based fraud detectors against camouflaged fraudsters. In: CIKM (2020)
- [4] McMahan, B., Moore, E., Ramage, D., Hampson, S., Arcas, B.A.y.: Communication-efficient learning of deep networks from decentralized data. In: Artificial Intelligence and Statistics (AISTATS) (2017)
- [5] Dai, E., Wang, S.: Say no to the discrimination: Learning fair graph neural networks with limited sensitive attribute information. In: WSDM (2021)
- [6] Yang, C., Liu, J., Yan, Y., Shi, C.: Fairsin: Achieving fairness in graph neural networks through sensitive information neutralization. In: AAAI (2024)
- [7] Lee, Y.-C., Shin, H., Kim, S.-W.: Disentangling, amplifying, and debiasing: Learning disentangled representations for fair graph neural networks. In: AAAI (2025)
- [8] Ezzeldin, Y.H., Yan, S., He, C., Ferrara, E., Avestimehr, A.S.: Fairfed: Enabling group fairness in federated learning. In: AAAI (2023)
- [9] Li, T., Sanjabi, M., Beirami, A., Smith, V.: Fair resource allocation in federated learning. In: International Conference on Learning Representations (ICLR) (2020)
- [10] Blanchard, P., El Mhamdi, E.M., Guerraoui, R., Stainer, J.: Machine learning with adversaries: Byzantine tolerant gradient descent. In: NeurIPS (2017)
- [11] Yin, D., Chen, Y., Kannan, R., Bartlett, P.: Byzantine-robust distributed learning: Towards optimal statistical rates. In: ICML (2018)
- [12] Kasyap, H., Fang, M., Liu, Z., Maple, C., Tripathy, S.: Fairness-constrained optimization attack in federated learning. arXiv preprint arXiv:2510.12143 (2025)
- [13] Agarwal, C., Lakkaraju, H., Zitnik, M.: Towards a unified framework for fair and stable graph representation learning. In: Uncertainty in Artificial Intelligence (UAI) (2021)
- [14] Dong, Y., Liu, N., Jalaian, B., Li, J.: Edits: Modeling and mitigating data bias for graph neural networks. In: The Web Conference (WWW) (2022)
- [15] Li, Z., Zhou, Y., *et al.*: Rethinking fair graph neural networks from re-balancing. In: KDD (2024)
- [16] Zhu, Y., Li, J., Chen, L., Zheng, Z.: One fits all: Learning fair graph neural

- networks for various sensitive attributes. In: KDD (2024)
- [17] Beam, S., *et al.*: Distribution-aware fairness for node representation learning on graphs. In: CIKM (2025)
  - [18] Corbucci, L., Monreale, A., *et al.*: Puffle: Balancing privacy, utility, and fairness in federated learning. In: ECML-PKDD (2024)
  - [19] Anonymous: Fedfact: Controllable global and local group fairness in federated learning. In: NeurIPS (2025)
  - [20] Meng, X., *et al.*: Fair federated graph learning with structure-aware group fairness. arXiv preprint arXiv:2310.12350 (2023)
  - [21] Zhou, Z., Yang, S., Zhao, F., Ren, X.: FairGFL: Privacy-preserving fairness-aware federated learning with overlapping subgraphs. arXiv preprint arXiv:2512.23235 (2025)
  - [22] FedGraph-Fair authors: Personalized federated graph learning with fairness via dynamic graphs and distributionally robust optimization. Information Sciences (2025)
  - [23] Wang, Z., Jin, Y.: Fairness-oriented vertical federated GNNs with incomplete sensitive attributes. Information Fusion **127**, 103871 (2026)
  - [24] Xue, G., Yu, Y.: Federated fairness-aware classification under differential privacy. arXiv preprint arXiv:2603.24392 (2026)
  - [25] McMahan, H.B., Ramage, D., Talwar, K., Zhang, L.: Learning differentially private recurrent language models. In: International Conference on Learning Representations (ICLR) (2018)
  - [26] Mironov, I.: Rényi differential privacy. In: IEEE Computer Security Foundations Symposium (CSF) (2017)
  - [27] Abadi, M., Chu, A., Goodfellow, I., McMahan, H.B., Mironov, I., Talwar, K., Zhang, L.: Deep learning with differential privacy. In: ACM CCS (2016)
  - [28] Ling, X., Fu, J., Chen, Z., *et al.*: FedFDP: Federated learning with fairness and differential privacy. arXiv preprint arXiv:2402.16028 (2024)
  - [29] El Mhamdi, E.M., Guerraoui, R., Rouault, S.: The hidden vulnerability of distributed learning in byzantium. In: ICML (2018)
  - [30] Cao, X., Fang, M., Liu, J., Gong, N.Z.: FLTrust: Byzantine-robust federated learning via trust bootstrapping. In: NDSS (2021)
  - [31] Bagdasaryan, E., Veit, A., Hua, Y., Estrin, D., Shmatikov, V.: How to backdoor

- federated learning. In: AISTATS (2020)
- [32] Baruch, G., Baruch, M., Goldberg, Y.: A little is enough: Circumventing defenses for distributed learning. In: NeurIPS (2019)
  - [33] Xie, C., Koyejo, O., Gupta, I.: Fall of empires: Breaking byzantine-tolerant sgd by inner product manipulation. In: Uncertainty in Artificial Intelligence (UAI) (2019)
  - [34] Meerza, S.I.A., Liu, J.: EAB-FL: Exacerbating algorithmic bias through model poisoning attacks in federated learning. In: IJCAI (2024)
  - [35] Gal, Y., Ghahramani, Z.: Dropout as a bayesian approximation: Representing model uncertainty in deep learning. In: ICML (2016)
  - [36] Hsu, T.-M.H., Qi, H., Brown, M.: Measuring the effects of non-identical data distribution for federated visual classification. arXiv preprint arXiv:1909.06335 (2019)
  - [37] Dwork, C., Roth, A.: The Algorithmic Foundations of Differential Privacy. Foundations and Trends in Theoretical Computer Science, ??? (2014)
  - [38] Jaggi, M.: Revisiting frank-wolfe: Projection-free sparse convex optimization. In: ICML (2013)

## Appendix A Hyperparameters and training protocol

Table A1 lists every hyperparameter used, with the value and the range we explored. All methods share the backbone, optimiser, and federated budget; method-specific hyperparameters (adversary learning rate for FairGNN,  $q$  for q-FedAvg,  $\beta$ -step for FairFed, post-processing for FDP-Fair) use the values recommended by their originating papers. Seeds are fixed ( $\{0, 1, 2\}$ ); differentially private runs are intrinsically stochastic, and we therefore report mean  $\pm$  standard deviation.

## Appendix B Dataset construction and sensitive attributes

**German**, **Credit**, **Bail** follow the standard NIFTY preprocessing: edges are symmetrised, features are  $z$ -score standardised, and the sensitive attribute is excluded from the node feature matrix (so the model cannot read it directly). Sensitive attributes are gender (German), age (Credit), and race (Bail); targets are good-credit, no-default, and (non-)recidivism respectively. **Pokec** uses region as the sensitive attribute and working-field as the target. **Elliptic** is the public Bitcoin transaction graph; illicit transactions are the positive class, and the unlabelled majority of nodes is retained



**Table A1:** Hyperparameters. “Explored” gives the sweep range where applicable.

| Hyperparameter             | Default                   | Explored                       |
|----------------------------|---------------------------|--------------------------------|
| GNN layers $L$             | 2                         | $\{2, 3\}$                     |
| Hidden dimension $d_h$     | 64                        | $\{32, 64, 128\}$              |
| Attention heads (FSER/GAT) | 4                         | $\{1, 4\}$                     |
| Dropout                    | 0.3                       | $\{0.2, 0.3, 0.5\}$            |
| Local optimiser            | AdamW                     | —                              |
| Local learning rate $\eta$ | 0.01                      | $\{10^{-3}, 10^{-2}\}$         |
| Weight decay               | $10^{-5}$                 | —                              |
| Local epochs $E$           | 2                         | $\{1, 2, 3\}$                  |
| Communication rounds $T$   | 40–60                     | per dataset                    |
| Clients $K$                | 3–10                      | $\{3, 5, 10, 20\}$             |
| Partition                  | Dirichlet( $\alpha=0.5$ ) | $\alpha \in \{0.1, 0.5, 1.0\}$ |
| Fairness weight $\lambda$  | 1.0                       | $\{0, 0.25, 0.5, 1, 2, 4\}$    |
| Fairness budget $\tau$     | 0.05                      | $\{0.01, 0.05, 0.1\}$          |
| DP $\epsilon$              | 8                         | $\{0.5, 1, 2, 4, 8, 16\}$      |
| DP $\delta$                | $10^{-5}$                 | —                              |
| DP clip $C$                | 1.0                       | —                              |
| Frank–Wolfe iters $T_{FW}$ | 20                        | —                              |
| Dual step $\eta_\mu$       | 0.1                       | —                              |
| MC-dropout samples         | 20                        | —                              |

for message passing but excluded from loss and metrics. As Elliptic carries no protected demographic attribute, we construct a transparent operational proxy: early-versus late-period transactions, split at the median of the 49 anonymised time steps, so that we study whether an AML model disproportionately flags one market period—a genuine temporal-drift concern in deployed compliance systems. We stress this is a subgroup-disparity analysis, not a protected-demographic claim.

Clients receive induced subgraphs under a Dirichlet( $\alpha$ ) label partition with a small per-client class floor that prevents degenerate single-class silos; each client inherits the global train/validation/test masks for its nodes, and the global model is evaluated on the pooled test nodes across all clients.

## Appendix C Reproducibility

All datasets are public and downloaded on demand into a local cache; all randomness is seeded; the privacy accountant, metrics, and aggregation rules are covered by an offline unit-test suite run in continuous integration; and every table and figure in this paper is regenerated from logged runs by a single script (`experiments/report.py`). The full experiment matrix—main comparison, ablation, privacy sweep, fairness–utility Pareto frontier, Byzantine-robustness study, scalability, and partition-sensitivity—is driven by a resumable runner (`experiments/run_matrix.py`). Code, config, and logs are released.